



From Fallbacks to Performance: Towards Better GlobalSel Performance on AArch64 NEON Platforms

Speaker: Madhur Amilkanthwar

In collaboration with:

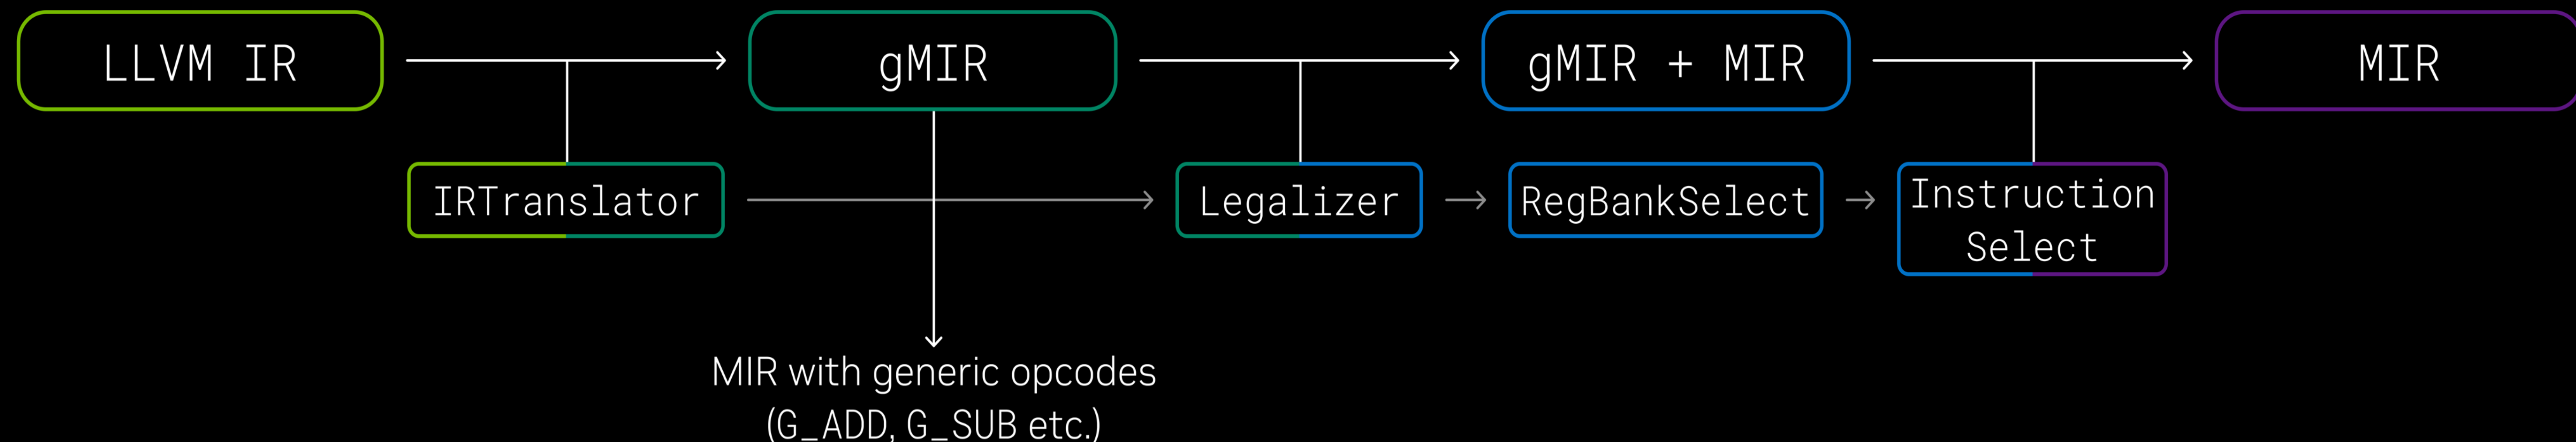
Dhruv Chawla, Tianyi Guan

Agenda

- GlobalSel – Quick Introduction and Status
- Our Contributions
- Results
- Next Steps

GlobalSel

- An alternate Instruction Selector that works on linear IR (GMIR, MIR) unlike DAG.
- Translates LLVM IR instructions to MIR in a phased manner.
- Promises compile-time improvements and offers platform to do function-wide optimizations.
- Default instruction selector at -O0 for AArch64 – fallbacks to SDAG for other optimization levels.
- **Our experience:**
 - TSVC, RajaPerf, SPEC 2017, LLVM Test-suite falls back to SDAG.



Goals

- Our long-term goal is to make GlobalSel the default instruction selector.
- We are taking an empirical approach and trying various known benchmarks with GlobalSel.
- A lot of work to do in supporting basic instructions and patterns.
- But... where do we start?
 - We start with something small and simple which allows us to get the ball rolling.
 - A small benchmark suite which can uncover issues in the framework.
 - Is there one out there which is widely known and small enough? Yes, **TSVC**!

TSVC and Scope

- TSVC – Test-Suite for Vectorizing Compilers
- Big file containing 152 **small** computation kernels focusing on loop optimizations.
- Kernel naming convention - **s<number>**

Our Scope:

- We focused on AArch64 Advanced SIMD a.k.a. NEON platform. Let's deal with SVE a bit later!
- Specifically targeting for `-mcpu=grace`.
- Compile with `-O3 -fglobal-isel`

Experimentation Setup:

- Isolated kernels to avoid i-cache and d-cache misses.
- Executed each kernel 5 times and taking minimum.

Status back in Feb 24:

One kernel falls back to SDAG – s317

TSVC s317 Kernel

```
real_t s317(struct args_t * func_args)
{
    // ...
    real_t q;
    for (int n1 = 0; n1 < 5*iterations; n1++) {
        q = (real_t)1.;
        for (int i = 0; i < LEN_1D/2; i++)
            q *= (real_t).99;
    }
    // ...
    return q;
}
```

- The loop was being vectorized with VF = 4, Unroll Factor = 2.



```
%vec.phi = phi <4 x float> [ <1, ...>, %ph ], [ %0, %vec.body ]
%vec.phi20 = phi <4 x float> [ <1, ...>, %ph ], [ %1, %vec.body ]
%0 = fmul fast <4 x float> %vec.phi, <0.99, ...>
%1 = fmul fast <4 x float> %vec.phi20, <0.99, ...>
```



```
%1 = phi <8 x float> [ %0, %vector.ph ], [ %3, %vector.body ]
%2 = call <8 x float> @llvm.vector.interleave2.v8f32(<4 x float> <0.99, ...>,
                                                    <4 x float> <0.99, ...>)
%3 = fmul fast <8 x float> %1, %2
```

TSVC s317 Kernel

Fixing the fallback

```
define <8 x i32> @test(<4 x i32> %a, <4 x i32> %b) {  
    %c = call <8 x i32> @llvm.vector.interleave2(<4 x i32> %a, <4 x i32> %b)  
    ret <8 x i32> %c  
}
```

SDAG



GISel



```
t2: v4i32,ch = CopyFromReg t0, Register:v4i32 %0  
t4: v4i32,ch = CopyFromReg t0, Register:v4i32 %1  
t5: v8i32 = concat_vectors t2, t4  
t7: v8i32 = vector_shuffle<0,4,1,5,2,6,3,7> t5,  
      undef:v8i32
```


```
bb.1 (%ir-block.0):  
    liveins: $q0, $q1  
    %0:_(<4 x s32>) = COPY $q0  
    %1:_(<4 x s32>) = COPY $q1  
    %2:_(<8 x s32>) = G_SHUFFLE_VECTOR %0:_(<4 x s32>),  
    %1:_, shufflemask(0, 4, 1, 5, 2, 6, 3, 7)
```

We added support for lowering the intrinsics to IRTranslator

[#85199](#) [GlobalSel] Add support for interleave and deinterleave intrinsics to IRTranslator

Performance Comparison

No fallback. Everything goes through Global ISel... but what about the performance?



TSVC s1279 Kernel

~65% slower on GlobalSel than SelectionDAG

```
real_t s1279(struct args_t * func_args)
{
    // ...
    for (int n1 = 0; n1 < iterations; n1++)
        for (int i = 0; i < LEN_1D; i++)
            if (a[i] < (real_t)0.)
                if (b[i] > a[i])
                    c[i] += d[i] * e[i];
    // ...
    return calc_checksum(__func__);
}
```

1

This select gets optimized into an AND after PreLegalizerCombiner, i.e.

```
%15:_(<4 x s1>) = G_FCMP floatpred(olt), %12:_(<4 x s32>), %13:_  
%17:_(<4 x s1>) = G_FCMP floatpred(ogt), %16:_(<4 x s32>), %12:_  
%72:_(<4 x s1>) = G_FREEZE %17:_  
%18:_(<4 x s1>) = G_AND %15:_, %72:_
```

3

The two if-conditions get merged into a single condition, with a select:

```
%0 = fcmp olt <4 x float> %wide.load, zeroinitializer  
%1 = fcmp ogt <4 x float> %wide.load39, %wide.load  
%2 = select <4 x i1> %0, <4 x i1> %1, <4 x i1> zeroinitializer
```

2

```
%164:_(s8) = G_TRUNC %180:_(s32)  
%162:_(s8) = G_TRUNC %181:_(s32)  
%169:_(s8) = G_FREEZE %168:_  
%167:_(s8) = G_FREEZE %166:_  
%165:_(s8) = G_FREEZE %164:_  
%163:_(s8) = G_FREEZE %162:_  
%187:_(s16) = G_ANYEXT %169:_(s8)  
%188:_(s16) = G_ANYEXT %167:_(s8)
```

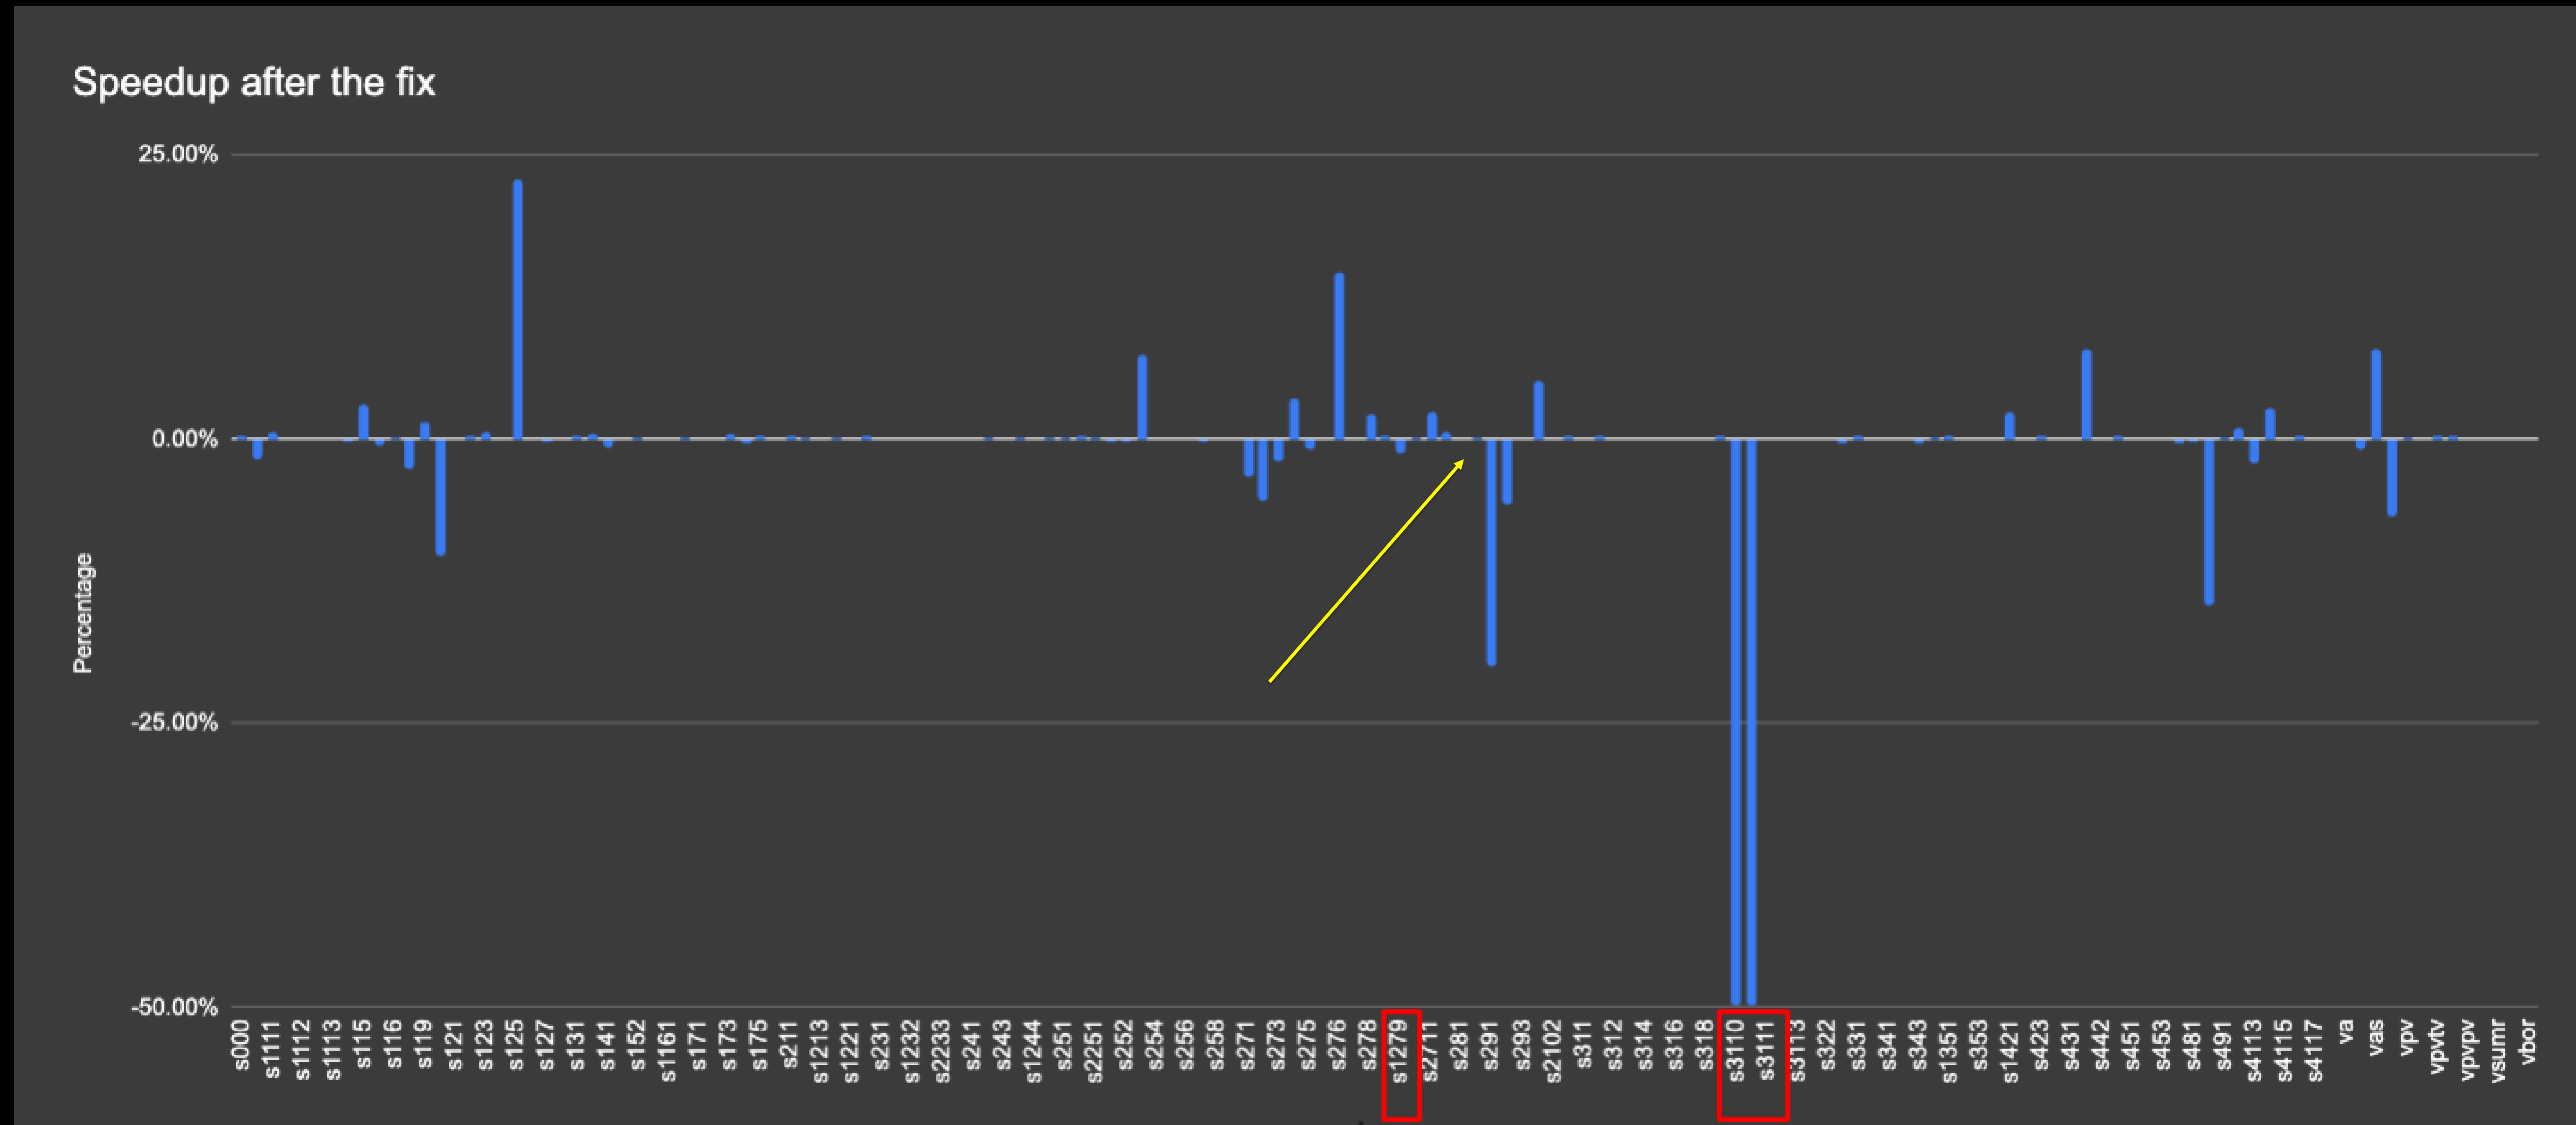
4

The Fixes

- **The fix adds rules to `AArch64LegalizerInfo.cpp` to avoid expansion.**
 - Issue: Legalizer was scalarizing any `G_FREEZE` not of the type `<2 x s64>`
 - Fix: legalize `G_FREEZE` like other vector operations
- **Solution:** [#88469](#) Avoid scalarizing `G_FREEZE` and `G_IMPLICIT_DEF`
- Other fixes:
 - [#89017](#) [AArch64] Fold `COPY(y:gpr, DUP(x:fpr, i)) -> UMOV(y:gpr, x:fpr, i)`
 - [#89023](#) [GISEL][CombinerHelper] Combine `op(trunc(x), trunc(y)) -> trunc(op(x, y))`

What do we gain?

- **Performance uplift:** Performance recovered, minor slowdown present.
- Collected at LLVM Commit: dbc1 (03 June 2024 - After our fixes)



TSVC s3110 Kernel

~50% slower on GlobalSel than SelectionDAG

```
real_t s3110(struct args_t * func_args)
{
    // ...
    real_t max;
    for (int n1 = 0; n1 < (...); n1++) {
        max = aa[0][0];
        for (int i = 0; i < LEN_2D; i++)
            for (int j = 0; j < LEN_2D; j++)
                if (aa[i][j] > max)
                    max = aa[i][j];
    }
    // ...
}
```

Source Code

LLVM IR

for.body:

%0 = load float, ptr @aa, align 64
br (...)

for.body.inner:

%3 = load float, ptr %<...>, align 4
%cmp = fcmp ogt float %3, %max
%max.2 = select i1 %cmp, float %3, float %max
br i1 (...)
}

TSVC s3110 Kernel

~50% slower on GlobalSel than SelectionDAG

Load of float
stored into GPR!




```
# Machine code for function s3110
# ...
bb.3.for.body:
    %0:gpr(s32) = G_LOAD %1:gpr(p0) ::
                (dereferenceable load (s32) from @aa, align 64)
```

```
bb.4.for.body.inner.preheader:
    %2:gpr(s32) = G_PHI %0:gpr(s32), %bb.3,
                  %8:fpr(s32), %bb.<...>

    G_BR (...)
```

```
bb.7.for.body.inner:
    %3:gpr(s32)  = G_PHI %2:gpr(s32), %bb.4,
                  %8:fpr(s32), %bb.7
    %6:fpr(s32) = G_LOAD %20:gpr(p0) ::
                  (load (s32) from %ir.<...>)
```

Unnecessary
copies from GPR
inserted



```
%4:fpr(s32) = COPY %3:gpr(s32)
%7:gpr(s32) = G_FCMP floatpred(ogt), %6:fpr(s32), %4:fpr
%5:fpr(s32) = COPY %3:gpr(s32)
%8:fpr(s32)  = G_SELECT %7:gpr(s32), %6:fpr, %5:fpr

    G_BR (...)
```

Issue:

RegBankSelect phase failed to assign optimal register banks to operands.

It did not analyze uses of G_LOAD beyond PHI nodes and assigned GPR bank which resulted in unnecessary copies.

The fix:

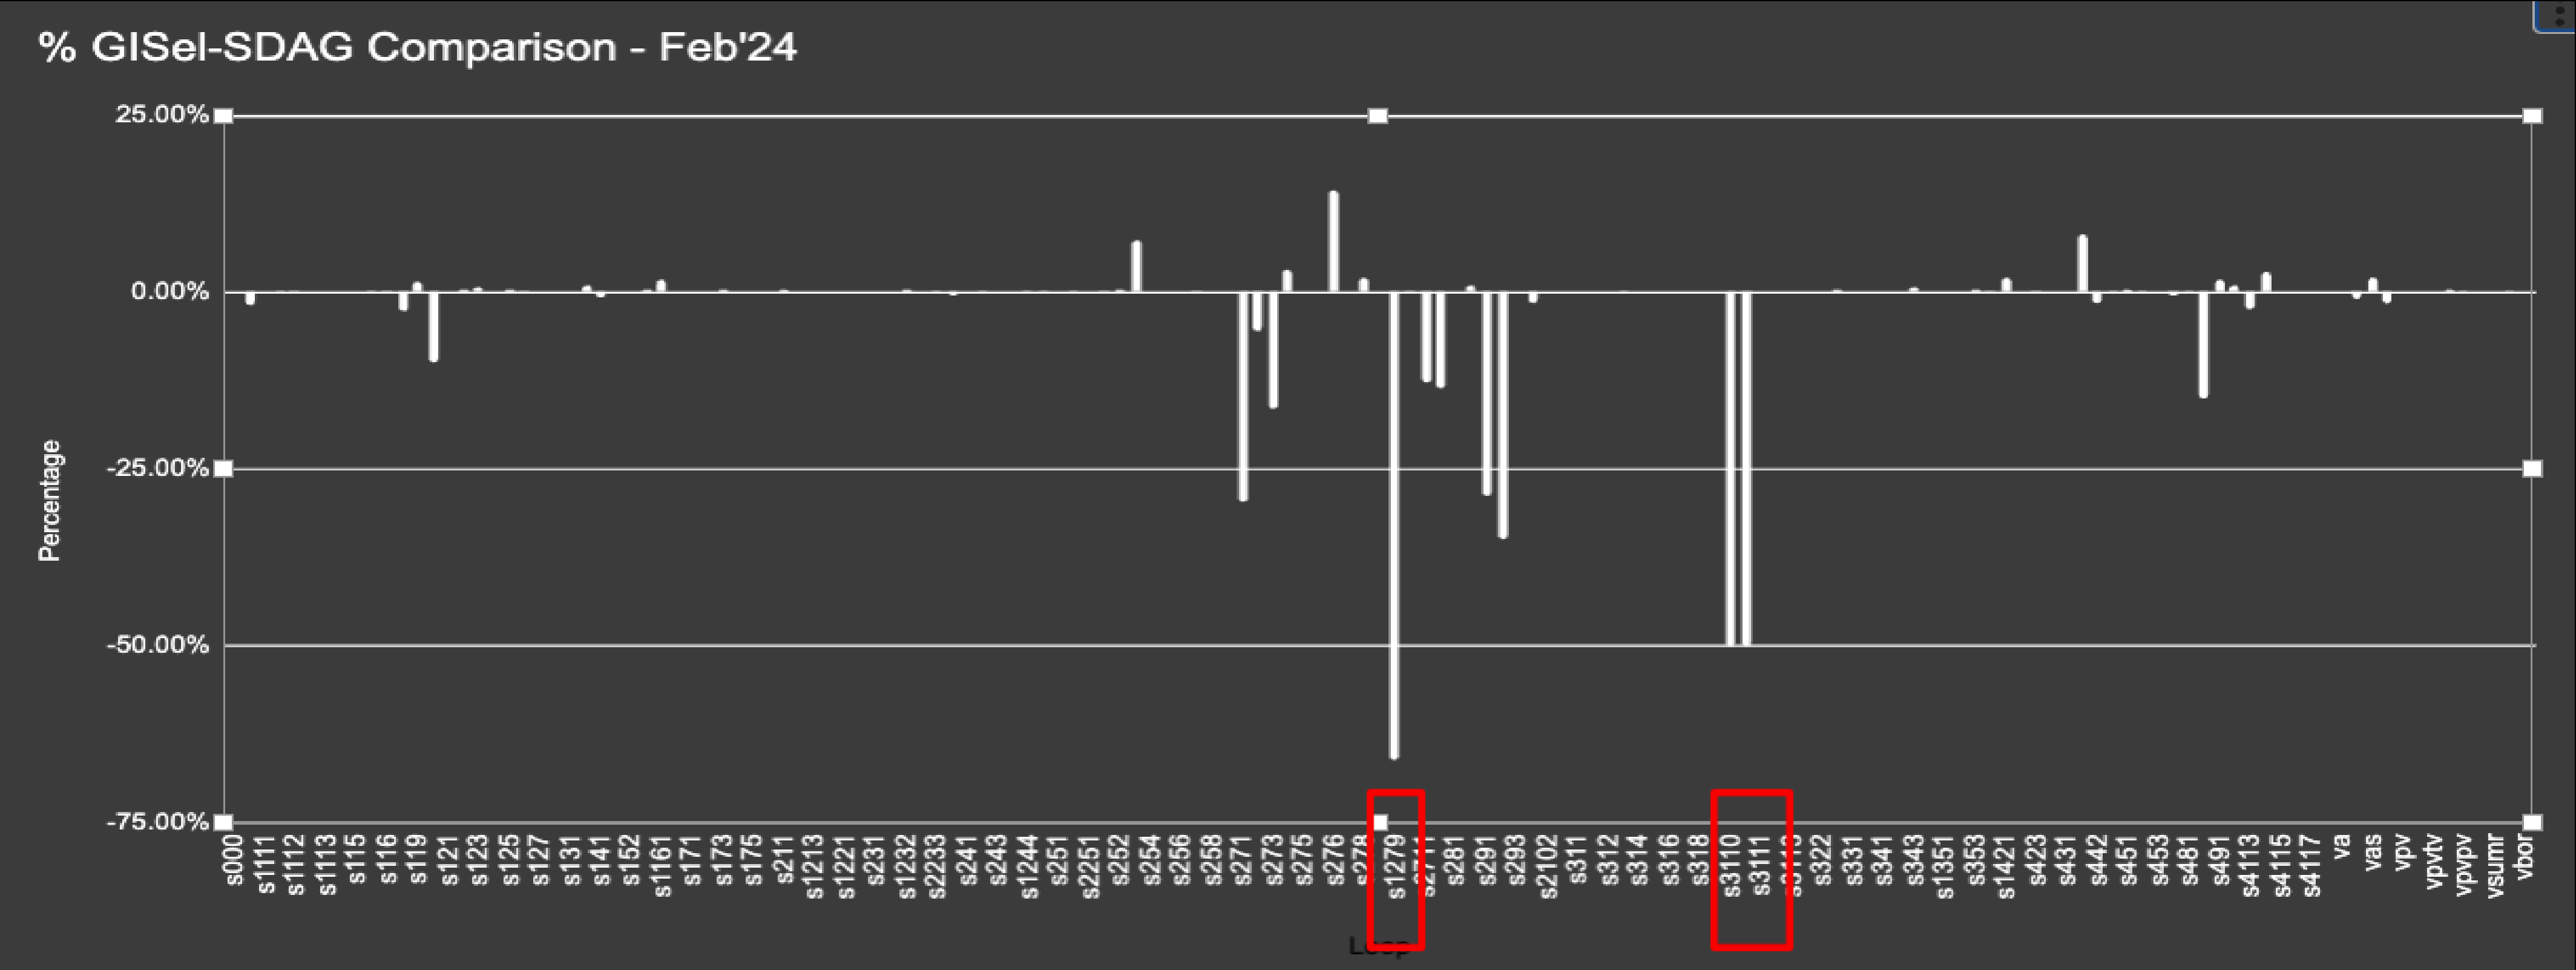
Analyze uses of G_LOAD in PHI node recursively.

If the def of PHI is used by instructions that only uses FP operands, then assign FPR to the def of G_LOAD.


The patch: [#94618](#)

Finally... where do we stand?

Before



After



Conclusions

Achievements:

- Fallback free TSVC at -O3 level for Advanced SIMD (Neon)
- Compile-time neutral performance parity of GlobalSel with Selection DAG for TSVC.

Future Work:

- Ongoing SVE support.
- Benchmark and workload analysis
 - SPEC 2017, RAJAPerf, LLVM Test-suite and further.
- GlobalSel by default is a long road
 - Patches welcome, come join us!

Current Status

Benchmark	# fallbacks	Notes
SPEC 2017, intrate	52	4 out of 9 benchmarks compile without fallbacks to SDAG.
RAJAPerf	6	G_EXT, G_AND and G_PTRTOINT
LLVM Test-suite	1029	Inline asm, G_EXT, G_FREEZE and many more are missing support.

our [patch](#) to support varargs in InstructionSelect phase reduced many fallbacks.



Thank you for attending!